LLM Abstention Study: Prompt Style Effects in Passage-Grounded Generation

Braeden Treutel
Department of Computer Science
Middle Tennessee State University
Murfreesboro, TN
Bmt5a@mtmail.mtsu.edu

1. Abstract

Large Language Models (LLMs) frequently generate confident but unsupported responses when confronted with questions they cannot answer, a phenomenon known as hallucination. Abstention, or the ability of a model to decline answering when sufficient evidence is unavailable, has been proposed as a mitigation strategy, but the mechanisms that drive reliable abstention behavior remain poorly understood. This study investigates whether explicit abstention instruction in a passage-grounded prompt is necessary for calibrated refusal behavior, or whether constraining a model to a provided passage is sufficient. Using LLaMA 3 8B Instruct on the SQuAD2.0 validation split, we compare two prompt configurations: an explicit condition that scripts a precise refusal phrase when the passage is insufficient, and an implicit condition that instructs the model to use only the passage without prescribing refusal language. Both conditions receive identical passage-question pairs drawn from SQuAD2.0, with 500 answerable and 500 unanswerable questions per condition. Results show that the explicit condition achieves a higher abstention rate on unanswerable questions (86.4% vs 55%) but shows greater over-abstention on answerable questions (78.4% answer attempt rate vs 96.6%). These findings suggest that explicit refusal scripting improves abstention on unanswerable questions but introduces a calibration cost on answerable ones, and that neither prompt strategy dominates across both dimensions.

2. Introduction

Large Language Models (LLMs) demonstrate strong performance across a wide range of natural language processing tasks, yet they frequently generate factually incorrect or unsupported responses, a phenomenon commonly referred to as hallucination. Unlike retrieval-based systems that surface existing documents, LLMs generate responses by predicting the most probable next token, optimizing likelihood rather than factual correctness. As a result, when confronted with questions they cannot reliably answer, LLMs often produce confident but unsupported responses rather than acknowledging uncertainty.

In high-stakes domains such as healthcare and legal assistance, this behavior carries significant consequences. A model that fabricates a medication dosage or misrepresents a legal standard does not only fail to help but also actively misleads. As LLM deployment expands into these domains, the ability to recognize and decline unanswerable questions

becomes as important as the ability to answer correctly. This capacity, known as abstention, has emerged as an important area of research in LLM safety and reliability.

This study compares two prompt configurations applied to the same passage-grounded pipeline: (1) an explicit condition that scripts a precise refusal phrase when the passage is insufficient, and (2) an implicit condition that constrains the model to the passage without prescribing refusal language. Both conditions receive identical passages and questions; only the prompt template differs. Because the same model and passages are used in both conditions, parametric knowledge is held constant, meaning any observed differences in abstention behavior can be attributed to prompt style rather than internal knowledge. This project aims to determine whether explicit refusal scripting is necessary for calibrated abstention or whether passage-grounding instruction alone is sufficient.

3. Background

LLMs are trained to predict the next token in a sequence, optimizing likelihood rather than factual correctness. As a result, they may generate fluent but unsupported responses when confronted with uncertainty [1]. Hallucination remains one of the most significant barriers to safe deployment of LLM systems, particularly in domains requiring factual reliability [1].

One mitigation strategy is abstention, which enables models to decline answering when they do not have the required context to correctly respond. Empirical studies show that LLMs frequently fail to abstain from unanswerable questions and instead generate confident but unsupported responses [3]. Previous work shows the potential of abstention for greater safety and reliability with limited context [3]. To address the lack of standardized evaluation, recent work introduces a black-box assessment framework and the Abstain-QA dataset, which explicitly separates answerable and unanswerable questions across domains and reasoning types [2]. While this work rigorously evaluates abstention in parametric LLMs, it does not examine how prompting strategies within a passage-grounded framework affect abstention ability.

Retrieval-augmented generation (RAG) integrates parametric knowledge with non-parametric external memory retrieved at inference time. By conditioning generation on external, verifiable knowledge, RAG aims to improve factual grounding and reduce hallucination [4]. While RAG is widely proposed as a hallucination mitigation strategy, existing evaluations primarily focus on answer accuracy rather than abstention behavior.

Recent work has begun analyzing hallucination evaluation under explicit abstention policies in RAG settings, demonstrating that common evaluation frameworks can produce very different judgements when refusal behavior is required [5]. However, this line of work evaluates the evaluators rather than isolating whether prompt-level instruction drives abstention. Therefore, it remains unclear whether explicit refusal scripting is necessary or whether constraining a model to the retrieved passage is sufficient to produce calibrated abstention. This project addresses that gap by systematically varying abstention instruction across two prompt conditions within a controlled passage-grounded pipeline.

4. Methods

4.1 Model and Experimental Design

LLaMA 3 8B Instruct was deployed locally using Ollama and queried via the Ollama Python API. Temperature was fixed at 0.0 across all conditions to ensure deterministic, reproducible outputs.

This study compares two prompt configurations applied to the same passage-grounded pipeline. Both conditions receive identical input of a Wikipedia passage paired with its corresponding question and differ only in the prompt template. Because the same model, passages, and questions are used in both conditions, any observed differences in abstention behavior can be attributed to prompt design rather than differences in parametric knowledge or retrieved content.

4.2 Prompt Conditions

The two conditions vary the degree of abstention instruction provided to the model. The explicit condition instructs the model to answer using only the provided passage and prescribes an exact refusal phrase to use when the passage is insufficient. The implicit condition instructs the model to use only the passage and to indicate when the answer is not present, without prescribing specific refusal language. Both prompts also include the passage and question in identical positions within the template. The full prompt templates are shown in Appendix A.

4.3 Dataset

The SQuAD2.0 validation split was used to evaluate abstention behavior under controlled answerable and unanswerable conditions. SQuAD2.0 pairs each question with a Wikipedia passage, and unanswerable questions were created by crowd workers who read each passage and deliberately authored plausible sounding questions that the passage does not answer. This design ensures that unanswerability is defined relative to the provided passage, which is consistent with the passage-grounded evaluation framework used in this study. The dataset was accessed via the HuggingFace Datasets library using the squad_v2 identifier and the validation split was loaded directly through the load_dataset API. The validation split contains 5,928 answerable and 5,945 unanswerable questions. A random sample of 500 answerable and 500 unanswerable questions was drawn using a fixed random seed, yielding 1,000 questions per condition and 2,000 total model queries. Each question was paired with exactly one passage; there was no retrieval across a document store, so cross-document contamination is not a concern.

4.4 Abstention Detection

Because SQuAD2.0 elicits free text, abstention was operationalized as any model output that explicitly declines to answer. Detection was performed using rule-based string matching against a predefined list of refusal phrases, including variants of “I don’t know”, “the passage does not”, “cannot be determined from”, and similar expressions. Outputs that triggered ambiguous phrases such as “it is unclear” or “without more context” were

flagged as edge cases and reviewed manually to ensure consistent classification. The full list of refusal/ambiguous phrases is shown in Appendix B. A response was classified as an abstention if and only if it contained a matching phrase. All other responses were classified as answer attempts.

4.5 Evaluation Metrics

Performance was assessed using three metrics. Abstention rate measures the proportion of unanswerable questions on which the model abstained and represents the primary measure of refusal behavior. Hallucination rate measures the proportion of unanswerable questions on which the model produced an answer rather than abstaining, and is defined as one minus the abstention rate. Answer attempt rate measures the proportion of answerable questions on which the model attempted an answer, and serves as a calibration check to detect over-abstention. A model that abstains on everything would score perfectly on abstention rate but near zero on answer attempt rate.

5. Results

5.1 Overview

Table 1 summarizes performance across both conditions on 500 answerable and 500 unanswerable questions drawn from the SQuAD2.0 validation split.

Condition	**Abstention Rate**	**Hallucination Rate**	**Answer Attempt Rate**
Explicit	86.4%	13.6%	78.4%
Implicit	55.0%	45.0%	96.6%

Table 1: Performance of explicit and implicit prompt conditions on 500 answerable and 500 unanswerable questions

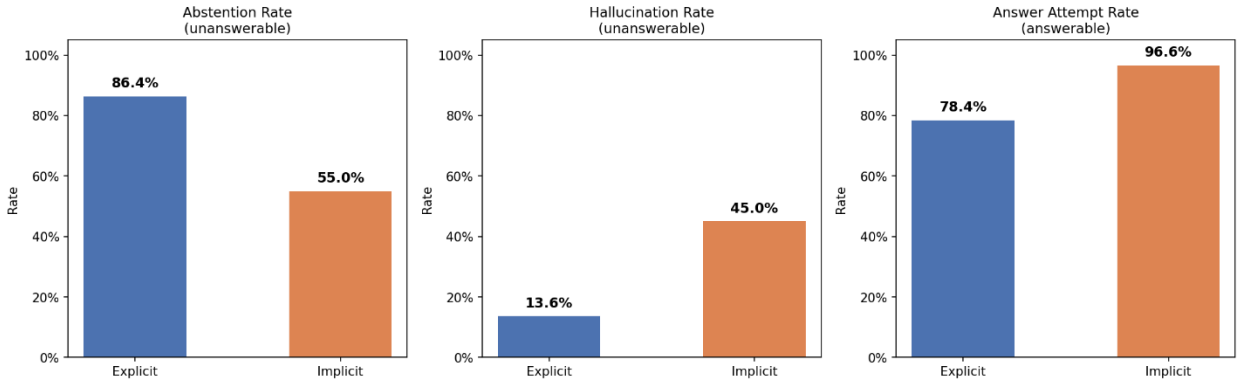

Figure 1: Abstention rate, hallucination rate, and answer attempt rate for explicit and implicit conditions

5.2 Abstention on Unanswerable Questions

The explicit condition abstained on 86.4% of unanswerable questions compared to 55.0% for the implicit condition, a difference of 31.4%. This result suggests that scripting a precise refusal phrase meaningfully increases the likelihood that a model will decline to answer when the passage does not contain sufficient information. Under the explicit prompt, the model has a concrete, low-effort output to produce when it cannot find an answer in the passage. Under the implicit prompt, the model must independently generate an appropriate refusal expression, which introduces more variability in whether or how it declines.

The following example illustrates a case where the two conditions diverge on an unanswerable question. Given the question "What P or was damaged during the 2008 tropical storm Fay?", which is a malformed question that asks about a specific entity the passage does not identify, the explicit condition correctly reasoned through the passage and abstained. The implicit condition, lacking a scripted refusal phrase, instead produced a factually grounded but non-responsive answer. It stated that parts of Jacksonville were under darkness for four days, without acknowledging that the specific entity asked about was not mentioned. This pattern, where the explicit condition correctly withholds and the implicit condition extracts an unsupported answer, was representative of the gap between the two conditions on unanswerable questions.

5.3 Answer Attempt Rate on Answerable Questions

While the explicit condition outperformed on unanswerable questions, it underperformed substantially on answerable ones. The implicit condition attempted to answer 96.6% of answerable questions compared to only 78.4% for the explicit condition, a difference of 18.2%. This indicates that the explicit prompt causes the model to over-abstain. It refuses to answer a meaningful proportion of questions that the passage does contain sufficient information to answer.

This behavior is consistent with the scripted refusal phrase acting as an easily reachable default. When the model encounters any uncertainty, even on questions that are technically answerable from the passage, the prescribed phrase provides a low-resistance output path. The implicit condition, lacking this scripted escape, is more likely to attempt an answer even under uncertainty, which benefits it on answerable questions but hurts it on unanswerable ones.

The Sky Q Hub question shows this concretely. Given the question "What is the name of Sky Q's broadband router?" against a passage that directly states, "a broadband router (Sky Q Hub)," the explicit condition incorrectly abstained. The implicit condition correctly extracted the answer, which suggests that the scripted refusal phrase in the explicit prompt can become a default output even when the passage contains a sufficient answer, particularly on questions requiring multi-step reasoning or questions where the relevant information is embedded within a list or parenthetical rather than stated as a standalone sentence.

5.4 The Calibration Tradeoff

Taken together, the results reveal a calibration tradeoff between the two conditions rather than a clear winner. The explicit condition is better calibrated on unanswerable questions but worse calibrated on answerable ones; the implicit condition shows the inverse pattern. Neither condition dominates across both dimensions simultaneously.

This tradeoff has practical implications for deployment. In a high-stakes domain where hallucination is the primary risk, such as medical or legal question answering, the explicit condition's lower hallucination rate may justify its higher rate of over-abstention. Conversely, in applications where failing to answer a valid question carries its own cost, such as a customer service or help desk system, the implicit condition's higher answer attempt rate may be preferable despite its higher hallucination rate.

6. Conclusion

This study investigated whether explicit abstention instruction is necessary for calibrated refusal behavior in passage-grounded generation, or whether constraining a model to a provided passage is sufficient. Using LLaMA 3 8B Instruct on a sample of 500 answerable and 500 unanswerable questions from the SQuAD2.0 validation split, we compared an explicit prompt condition that scripts a precise refusal phrase against an implicit condition that instructs the model to use only the passage without prescribing refusal language.

Results show that explicit abstention instruction substantially improves abstention on unanswerable questions (86.4% vs 55.0%) but introduces a meaningful calibration cost on answerable ones (78.4% vs. 96.6% answer attempt rate). Neither condition dominates across both dimensions. These findings suggest that the scripted refusal phrase in the explicit condition functions as a low-resistance default output that the model reaches for even on questions that the passage can answer, particularly when the reasoning required is non-trivial. The implicit condition, lacking this escape route, is better calibrated on answerable questions but less likely to abstain on unanswerable ones.

Future work should examine whether the over-abstention observed in the explicit condition can be reduced through refining the prompt. For instance, adding an instruction that explicitly discourages unnecessary refusals on answerable questions, or combining a scripted refusal phrase with a confidence-based condition, are both good directions for future work. It would also be valuable to evaluate this tradeoff across different models and model sizes to determine whether the calibration cost of explicit instruction is specific to LLaMA 3 8B or a more general phenomenon.

7. References

[1] Rawte, V., Sheth, A., & Das, A. (2023). *A Survey of Hallucination in Large Foundation Models*. arXiv. https://doi.org/10.48550/ARXIV.2309.05922

[2] Madhusudhan, N., Madhusudhan, S. T., Yadav, V., & Hashemi, M. (2025). Do LLMs Know When to NOT Answer? Investigating Abstention Abilities of Large Language Models. In O. Rambow, L. Wanner, M. Apidianaki, H. Al-Khalifa, B. D. Eugenio, & S. Schockaert (Eds.), *Proceedings of the 31st International Conference on Computational Linguistics* (pp. 9329–9345). Association for Computational Linguistics. https://aclanthology.org/2025.coling-main.627/

[3] Wen, B., Yao, J., Feng, S., Xu, C., Tsvetkov, Y., Howe, B., & Wang, L. L. (2025). Know Your Limits: A Survey of Abstention in Large Language Models. *Transactions of the Association for Computational Linguistics*, *13*, 529–556. https://doi.org/10.1162/tacl_a_00754

[4] Ma, Y., Nie, H., Chen, C., Zhang, J., Jiang, J., Wang, B., & Xia, Y. (2025). A Survey of Retrieval-Augmented Generation (RAG) for Large Language Models. *2025 International Conference on Trustworthy Big Data and Artificial Intelligence (ICTBAI)*, 7–13. https://doi.org/10.1109/ICTBAI68361.2025.00008

[5] Lamhot Siagian. (2026). *Benchmarking Hallucination Evaluation for RAG Under an Abstention Policy: A Controlled 30-Query Study with RAGAS, DeepEval, and LLM-as-Judge*. https://doi.org/10.13140/RG.2.2.23948.78726

8. Appendix

Appendix A: Full prompt templates

The following prompt templates were used in all experimental conditions. The `{context}` and `{question}` placeholders are filled at inference time with the SQuAD2.0 passage and question respectively.

```
PROMPT_EXPLICIT = """Answer the question using only the
information in the passage below.
If the passage does not contain enough information to answer the
question, you must respond with exactly: "The passage does not
contain enough information to answer this question."
Do not use any outside knowledge.

Passage:
{context}

Question: {question}"""

PROMPT_IMPLICIT = """Answer the question using only the
information in the passage below.
If the passage does not contain the answer, say so.
Do not use any outside knowledge.

Passage:
{context}

Question: {question}"""
```

Appendix B: List of refusal/ambiguous phrases

The following phrase lists were used for rule-based abstention detection. Outputs matching any phrase in the abstention list were classified as abstentions. Outputs matching any phrase in the edge case list were flagged for manual review.

```
# phrases that clearly indicate abstention
ABSTENTION_PHRASES = [
    # generic uncertainty
    "i don't know",
    "i do not know",
    "i'm not sure",
    "i am not sure",
    "i cannot answer",
    "i can't answer",
    "i don't have enough information",
    "i do not have enough information",
    "i don't have information",
```

```
    "not enough information",
    "insufficient information",
    "no information",
    # passage-grounded style refusals
    "the passage does not",
    "the passage doesn't",
    "the context does not",
    "the context doesn't",
    "not mentioned in the passage",
    "not stated in the passage",
    "not provided in the passage",
    "not found in the passage",
    "not mentioned in the context",
    "cannot be determined from",
    "cannot be found in",
    "no answer can be found",
    "there is no information",
    "the text does not",
    "based on the passage, i cannot",
    "based on the provided",
    # direct refusals
    "cannot provide an answer",
    "unable to answer",
    "this question cannot be answered",
    "there is no answer",
    "no answer is provided",
]

# phrases that are more ambiguous
# flag these for manual review
EDGE_CASE_PHRASES = [
    "it is unclear",
    "it's unclear",
    "it is not clear",
    "unclear from",
    "it depends",
    "may or may not",
    "i would need more",
    "without more context",
    "without additional",
]
```

Appendix C: Coding environment

The following describes the environment used to run the experiment. The setup is sufficient to recreate the coding environment.

Hardware

- CPU: Intel Core i9-9900K @ 3.60 GHz (8 cores, 16 logical processors)
- RAM: 128 GB
- GPU: None, performed on CPU
- OS: Windows 11 Pro with Ubuntu via Windows Subsystem for Linux (WSL)

Software

- Python 3.12.3
- Ollama 0.18.0
- Model: LLaMA 3 8B (llama3:8b), accessed locally via Ollama

Environment Setup

- The project repository can be cloned from GitHub using `'git clone https://github.com/braeden512/prompt-abstention-eval.git'`.
- Once cloned, navigate to the project root and create a Python virtual environment using `'python3 -m venv venv'`, then activate it with `'source venv/bin/activate'` on Linux, or `'venv\Scripts\activate'` on Windows.
- Navigate to the `src` directory using `'cd src'` and install dependencies with `'pip install -r requirements.txt'`.
- Ensure Ollama is installed using `'sudo snap install ollama'`, after which the model can be pulled with `'ollama pull llama3:8b'`.
- To run the experiment, navigate to the `src` directory and execute `'python run_experiment.py'` for the full run. After the experiment is completed, run `'python analyze_results.py'` to generate plots and summary statistics. Results are saved to the `results` directory in the project root.

A full list of installed packages is available in `'requirements.txt'`, located in the project source code repository.